

Scala/Spark

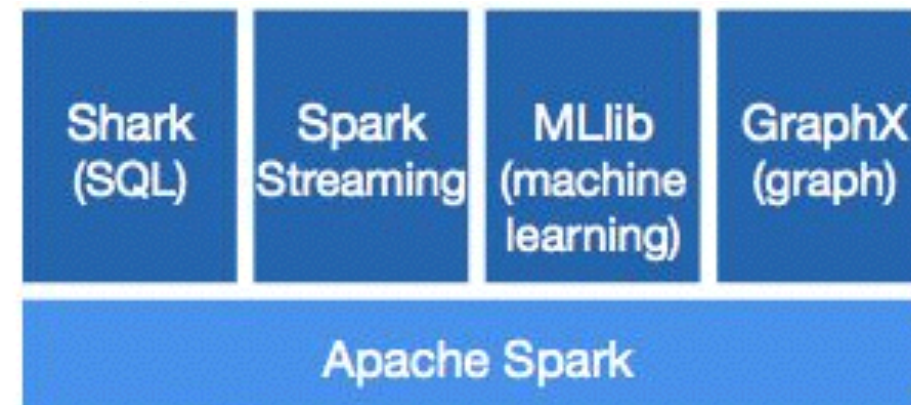


Wikipedia :

Apache Spark is an [open-source](#) cluster computing framework originally developed in the AMPLab at [UC Berkeley](#). In contrast to [Hadoop](#)'s two-stage disk-based [MapReduce](#) paradigm, Spark's in-memory primitives provide performance up to 100 times faster for certain applications.^[1] By allowing user programs to load data into a cluster's memory and query it repeatedly, Spark is well-suited to machine learning algorithms.^[2] Spark requires a cluster manager and a distributed storage system. For cluster management, Spark supports standalone (native Spark cluster), [Hadoop YARN](#), or [Apache Mesos](#).^[3] For distributed storage, Spark can interface with a wide variety, including [Hadoop Distributed File System \(HDFS\)](#),^[4] [Cassandra](#),^[5] [OpenStack Swift](#), and [Amazon S3](#). Spark also supports a pseudo-distributed local mode, usually used only for development or testing purposes, where distributed storage is not required and the local file system can be used instead; in this scenario, Spark is running on a single machine with one executor per [CPU core](#).

Project Components


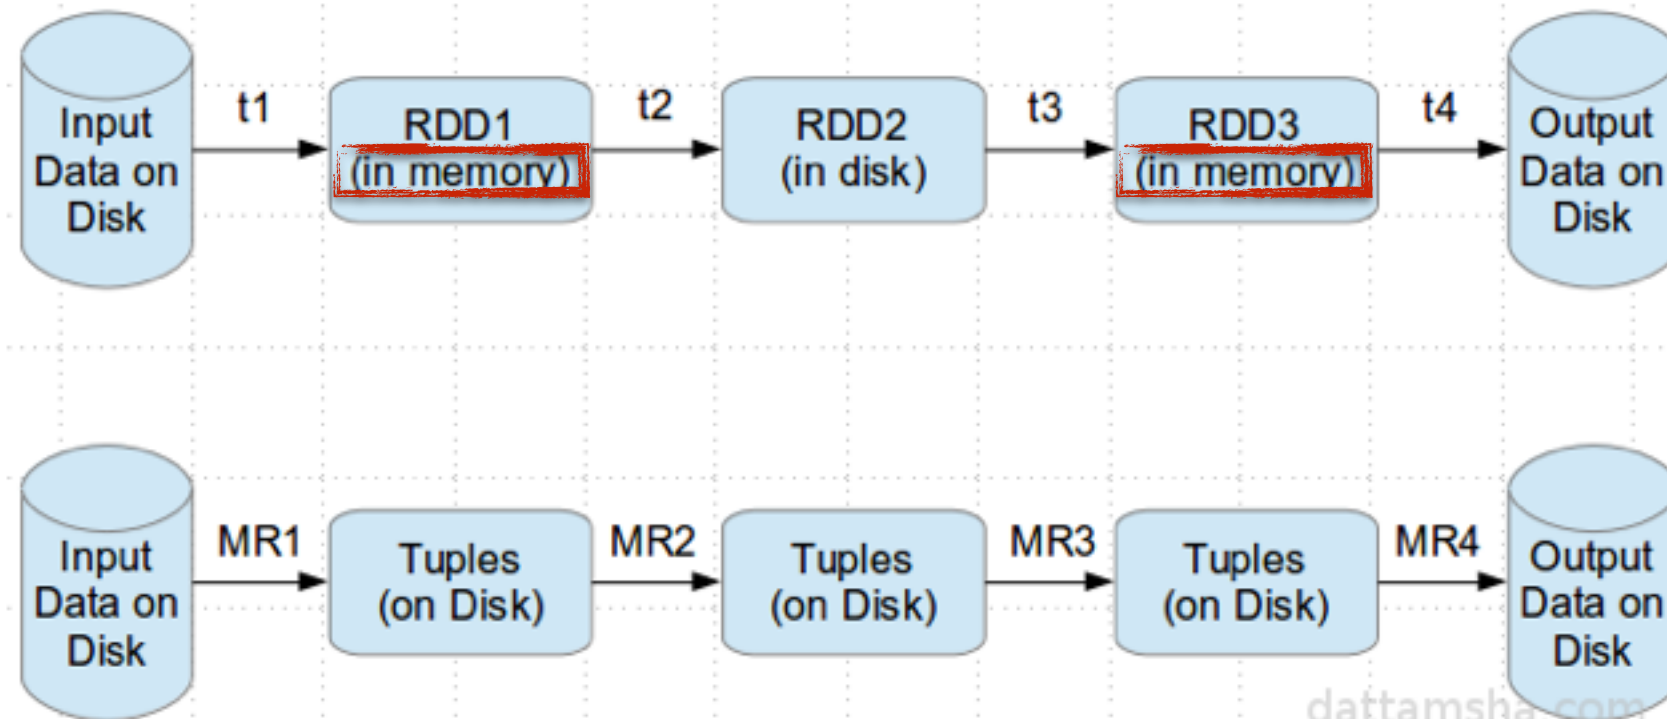
1. Spark Core and Resilient Distributed Datasets (RDDs)
2. Spark SQL
3. Spark Streaming
4. MLlib Machine Learning Library
5. GraphX



Features

- Java, Scala, Python, and R APIs.
- Scalability to over 8000 nodes in production.^[11]
- Ability to cache datasets in memory for interactive data analysis: extract a working set, cache it, query it repeatedly.
- Interactive command line interface (in Scala or Python) for low-latency [horizontally scalable](#) data exploration.
- Higher level library for stream processing^[12], through Spark Streaming.
- Support for structured and relational query processing (SQL), through Spark SQL.
- Higher level libraries for machine learning and graph processing.

Hadoop vs Spark

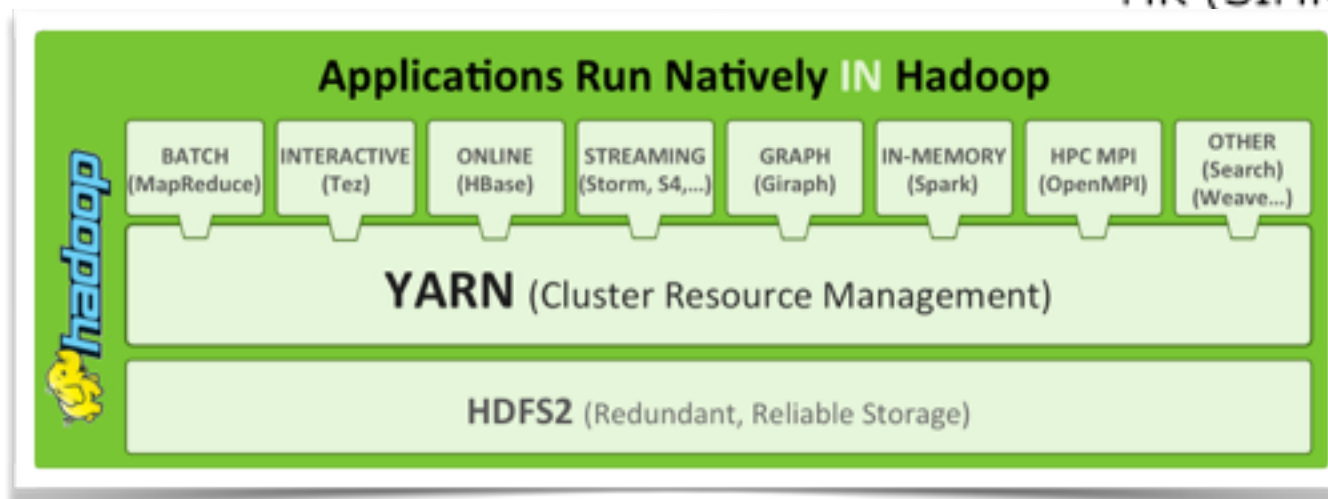
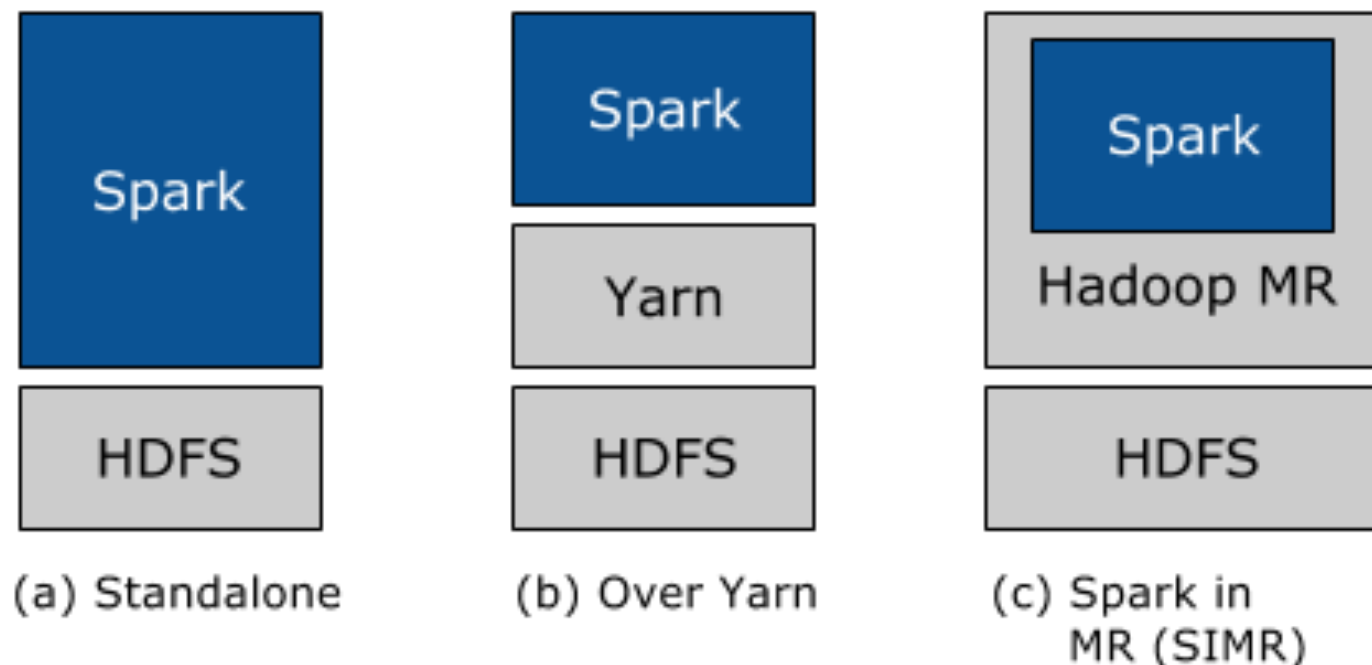


→ Move data through disk and network (HDFS)




→ User can cache data in memory

Spark and Hadoop : Working Together



Spark In MapReduce ([SIMR](#)): For the Hadoop users that are not running YARN yet, another option, in addition to the standalone deployment, is to use SIMR to launch Spark jobs inside MapReduce. With SIMR, users can start experimenting with Spark and use its shell within a couple of minutes after downloading it! This tremendously lowers the barrier of deployment, and lets virtually everyone play with Spark.



Scala

History

The design of Scala started in 2001 at the [École Polytechnique Fédérale de Lausanne](#) (EPFL) by [Martin Odersky](#), following on from work on Funnel, a programming language combining ideas from functional programming and [Petri nets](#).^[10] Odersky had previously worked on [Generic Java](#) and [javac](#), Sun's Java compiler.^[10]

After an internal release in late 2003, Scala was released publicly in early 2004 on the [Java](#) platform,^[11] and on the [.NET platform](#) in June 2004.^{[6][10][12]} A second version (v2.0) followed in March 2006.^[6] The .NET support was officially dropped in 2012.^[13]

Although Scala had extensive support for functional programming from the beginning, Java remained a purely object oriented language until the introduction of [lambda expressions](#) with [Java 8](#) in 2014.

- * In Scala, the compiler is incredibly smart, so this avoids the developer needing to specify explicitly those things that the compiler can infer.
- * Scala is ported on Spark (java libraries can be used)

Build Scala project with SBT



<http://www.scala-sbt.org/0.13/tutorial/index.html>

Scala crash course

<https://www.sics.se/~amir/files/download/dic/scala.pdf>



Examples Scala

- **Declare a list of integers as a variable called “myNumbers”:**

```
scala> val myNumbers = List(1, 2, 5, 4, 7, 3)
myNumbers: List[Int] = List(1, 2, 5, 4, 7, 3)
```

- **Declare a function, cube, that computes the cube (third power) of an Int.**

```
scala> def cube(a: Int): Int = a * a * a
cube: (a: Int)Int
```

- **Apply the function to myNumbers using the map function.**

```
scala> myNumbers.map(x => cube(x))
res: List[Int] = List(1, 8, 125, 64, 343, 27)
// Scala also provides some shorthand ways of writing this:
// myNumbers.map(cube(_))
// myNumbers.map(cube)
```

- **Then also try writing the function inline in a map call, using closure notation.**

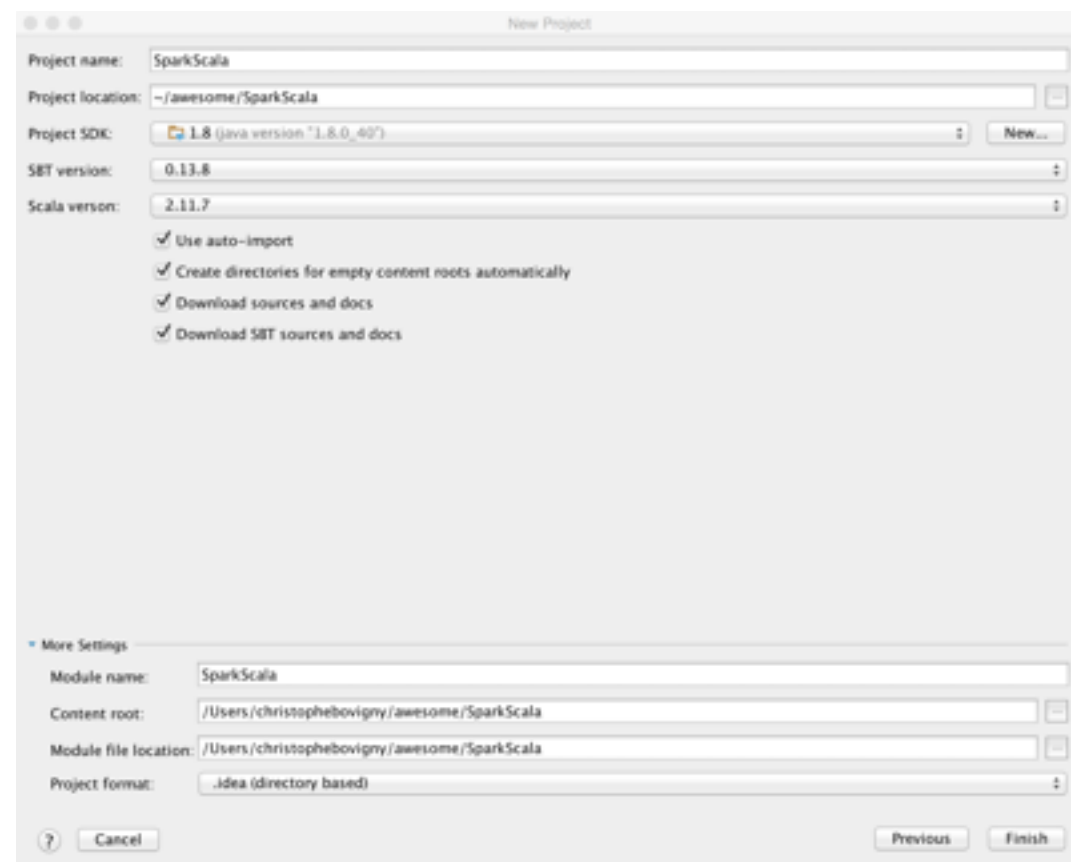
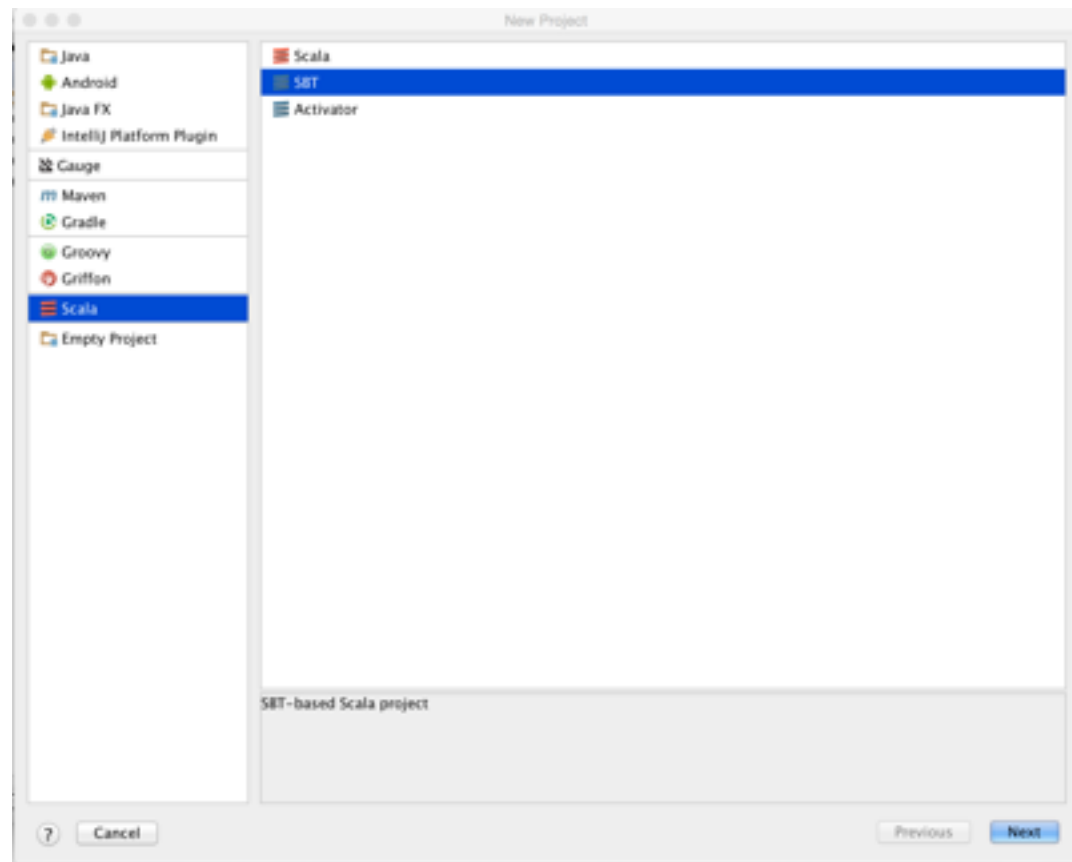
```
scala> myNumbers.map{x => x * x * x}
res: List[Int] = List(1, 8, 125, 64, 343, 27)
```

- **Define a factorial function that computes $n! = 1 * 2 * \dots * n$ given input n . You can use either a loop or recursion, in our solution we use recursion (see steps 5-7 of [First Steps to Scala](#)). Then compute the sum of factorials in myNumbers.**

```
scala> def factorial(n: Int): Int = if (n == 0) 1 else n * factorial(n - 1) // From http://bit.ly/b2sVKI
factorial: (Int)Int
scala> myNumbers.map(factorial).sum
res: Int = 5193
```



Setup a new Project





Add dependencies

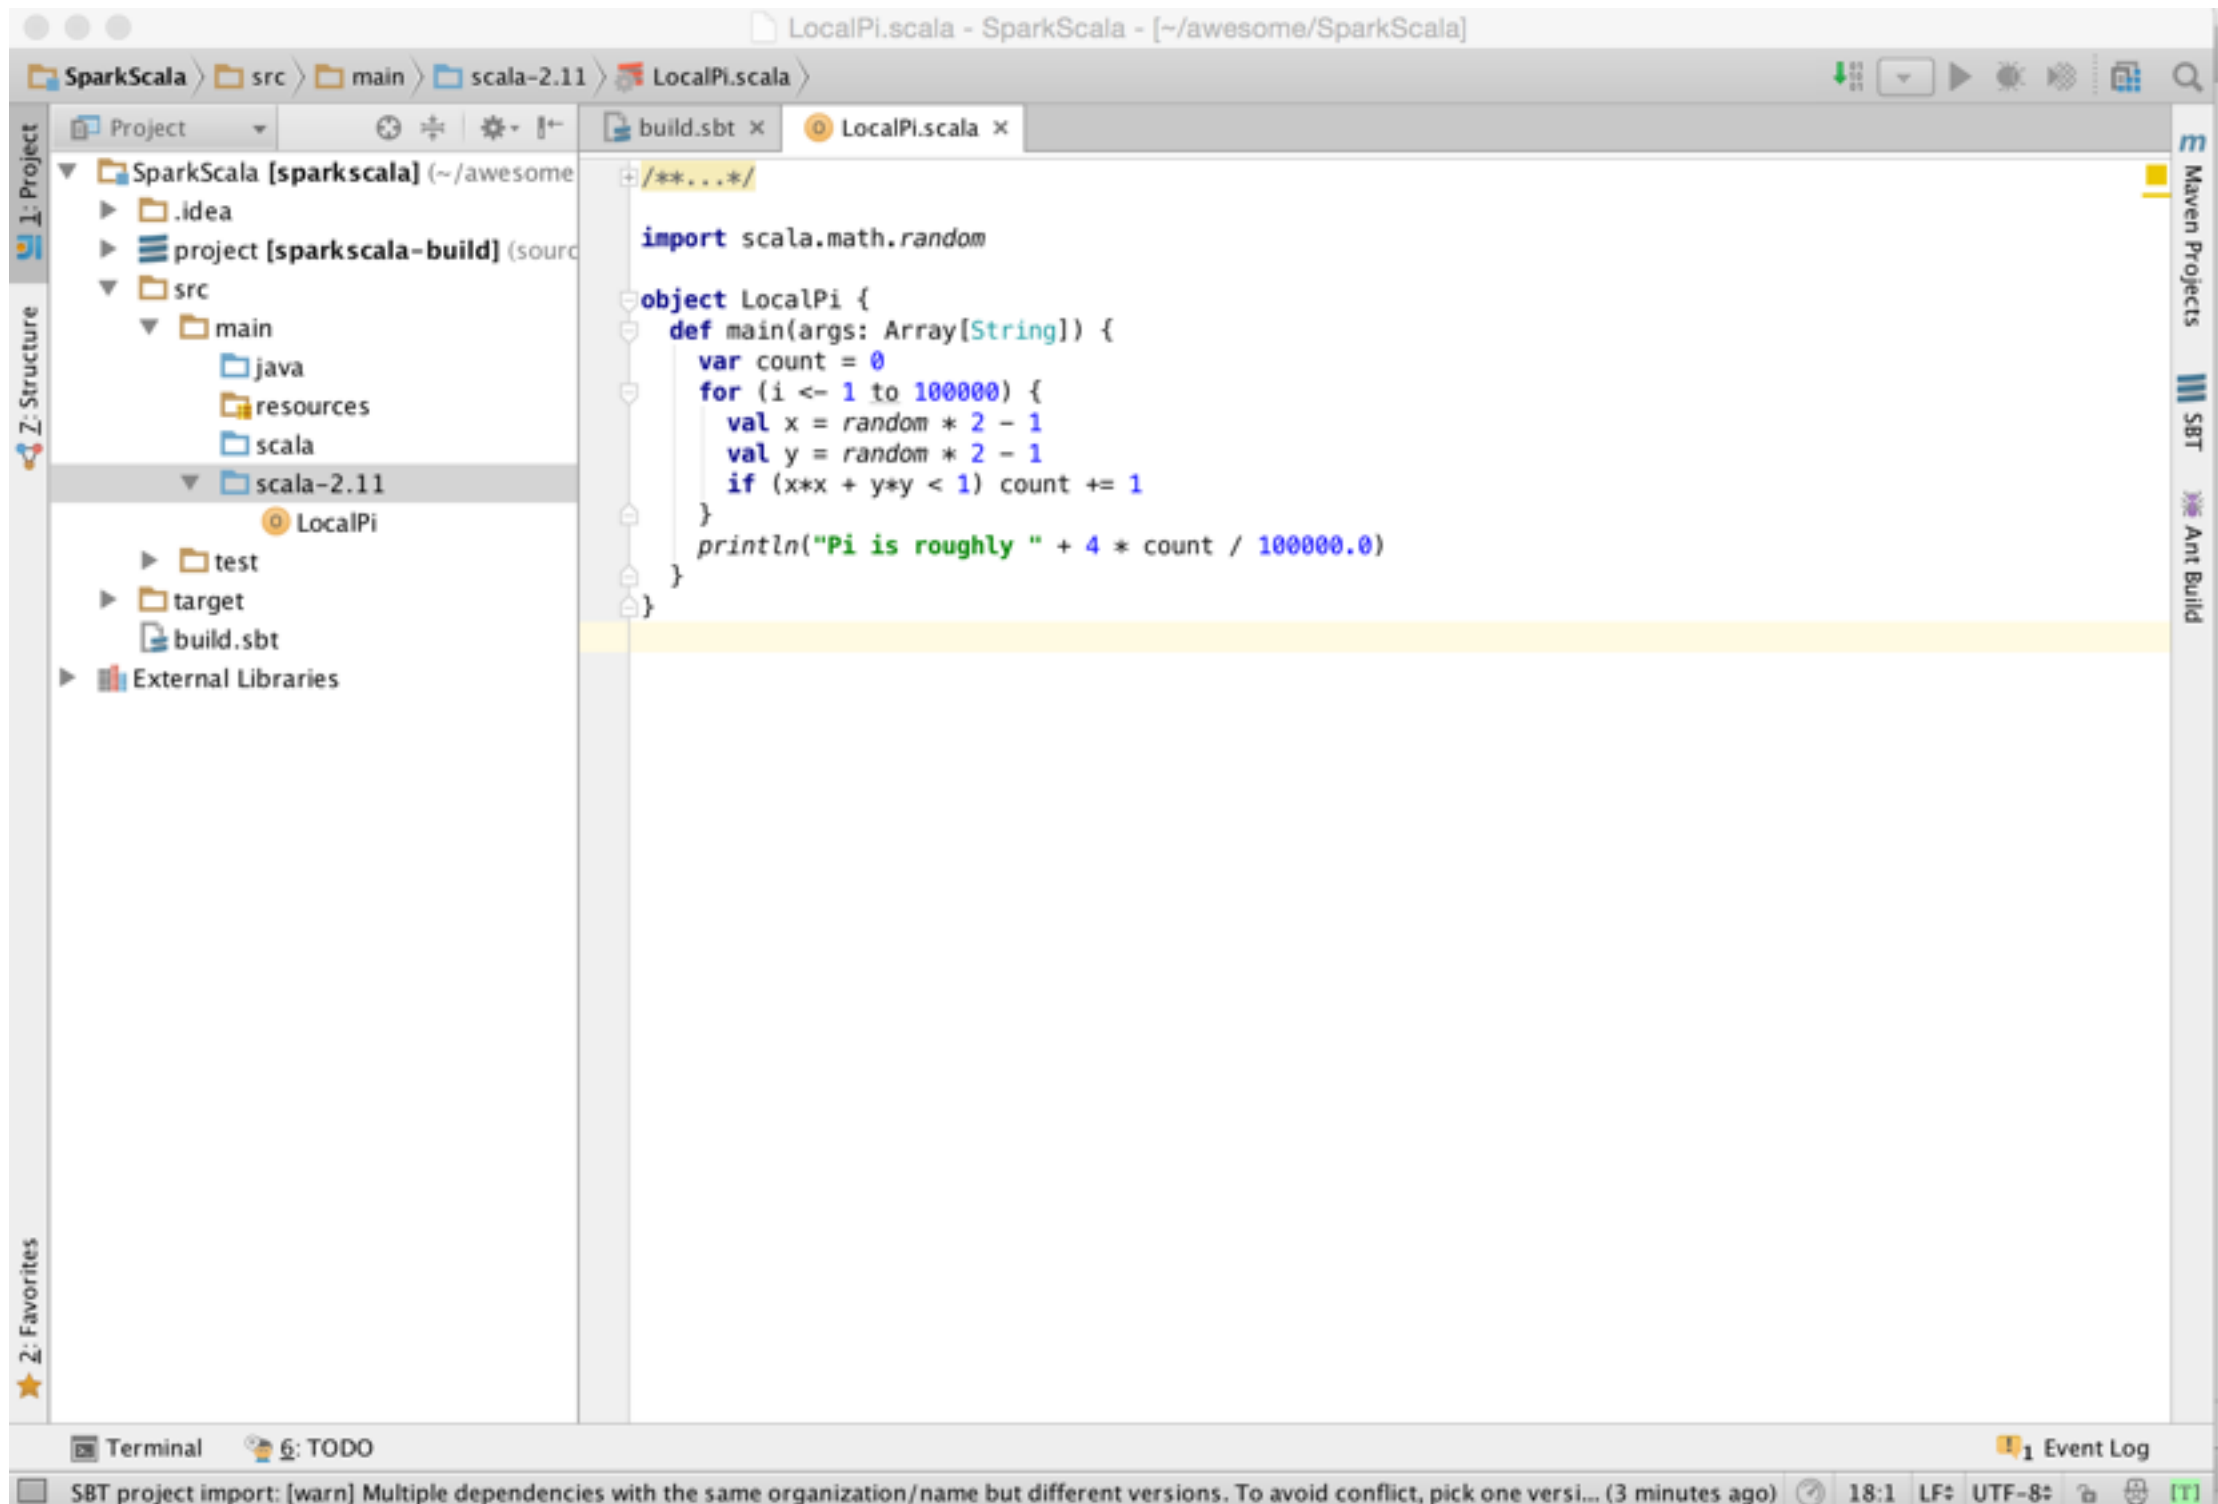
```
build.sbt - SparkScala - [~/awesome/SparkScala]

name := "SparkScala"
version := "1.0"
scalaVersion := "2.11.7"

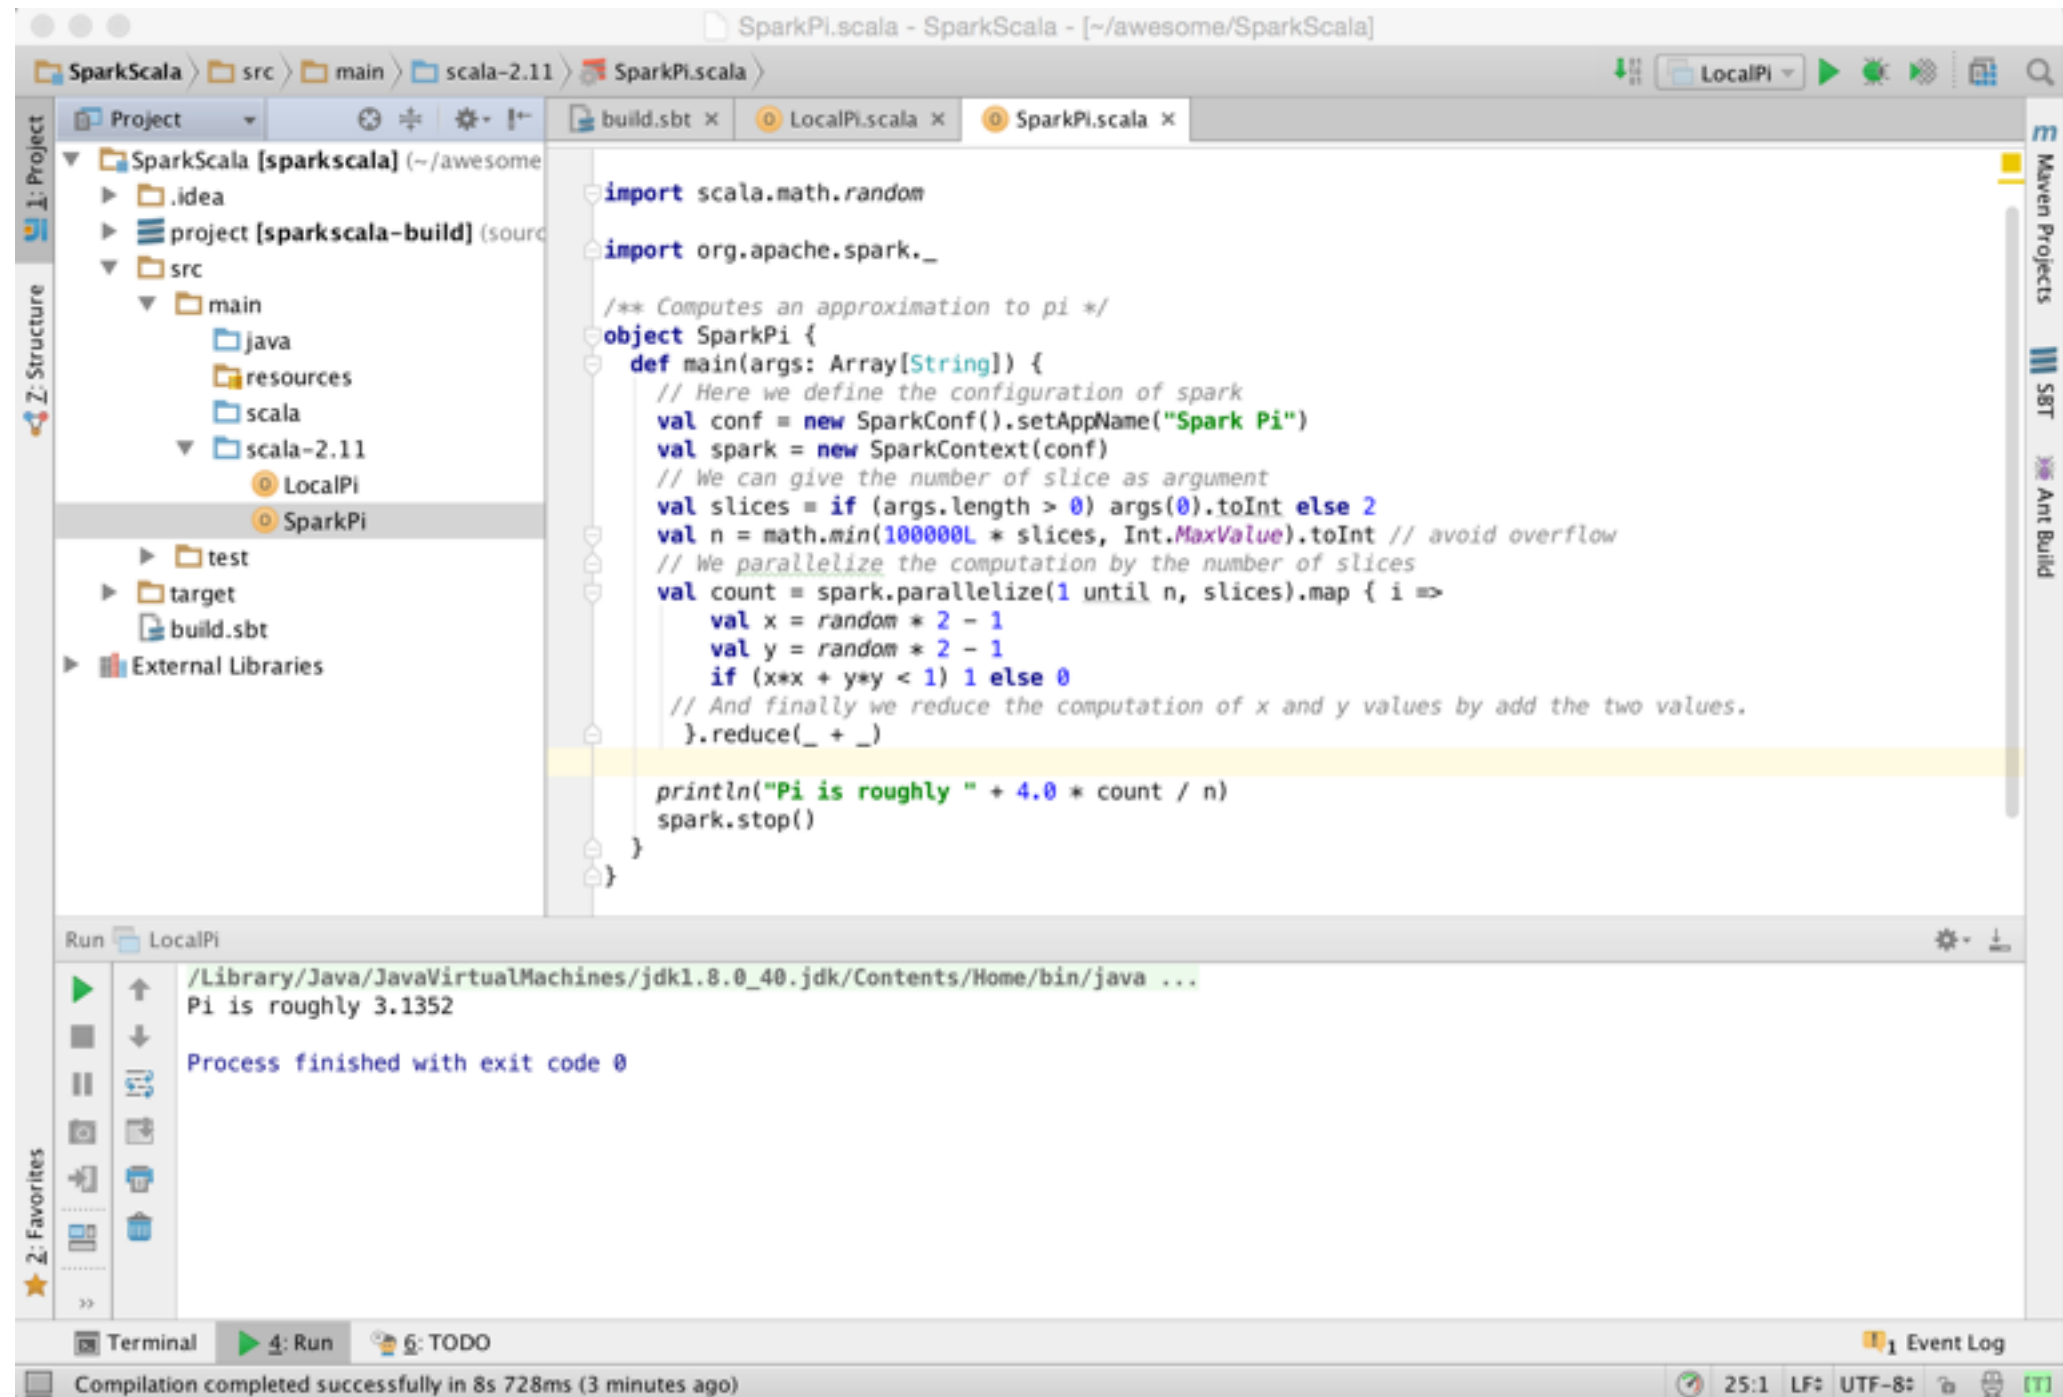
libraryDependencies ++= Seq(
  "org.apache.spark" %% "spark-streaming" % "1.4.1",
  "org.apache.spark" %% "spark-streaming-kafka" % "1.4.1",
  "javax.servlet" % "javax.servlet-api" % "3.0.1"
)
```

The screenshot shows an IDE window titled 'SparkScala' with a file named 'build.sbt' open. The left sidebar shows the project structure with folders like '.idea', 'project', 'src', 'target', and 'build.sbt'. The main editor area displays the Scala code for the build file. The code defines the project name as 'SparkScala', the version as '1.0', and the Scala version as '2.11.7'. It also defines the library dependencies as a sequence of three dependencies: 'org.apache.spark' %% 'spark-streaming' % '1.4.1', 'org.apache.spark' %% 'spark-streaming-kafka' % '1.4.1', and 'javax.servlet' % 'javax.servlet-api' % '3.0.1'. The bottom status bar shows 'SBT: [info] Loading project definition from /Users/christophebovigny/awesome/SparkSc...'.

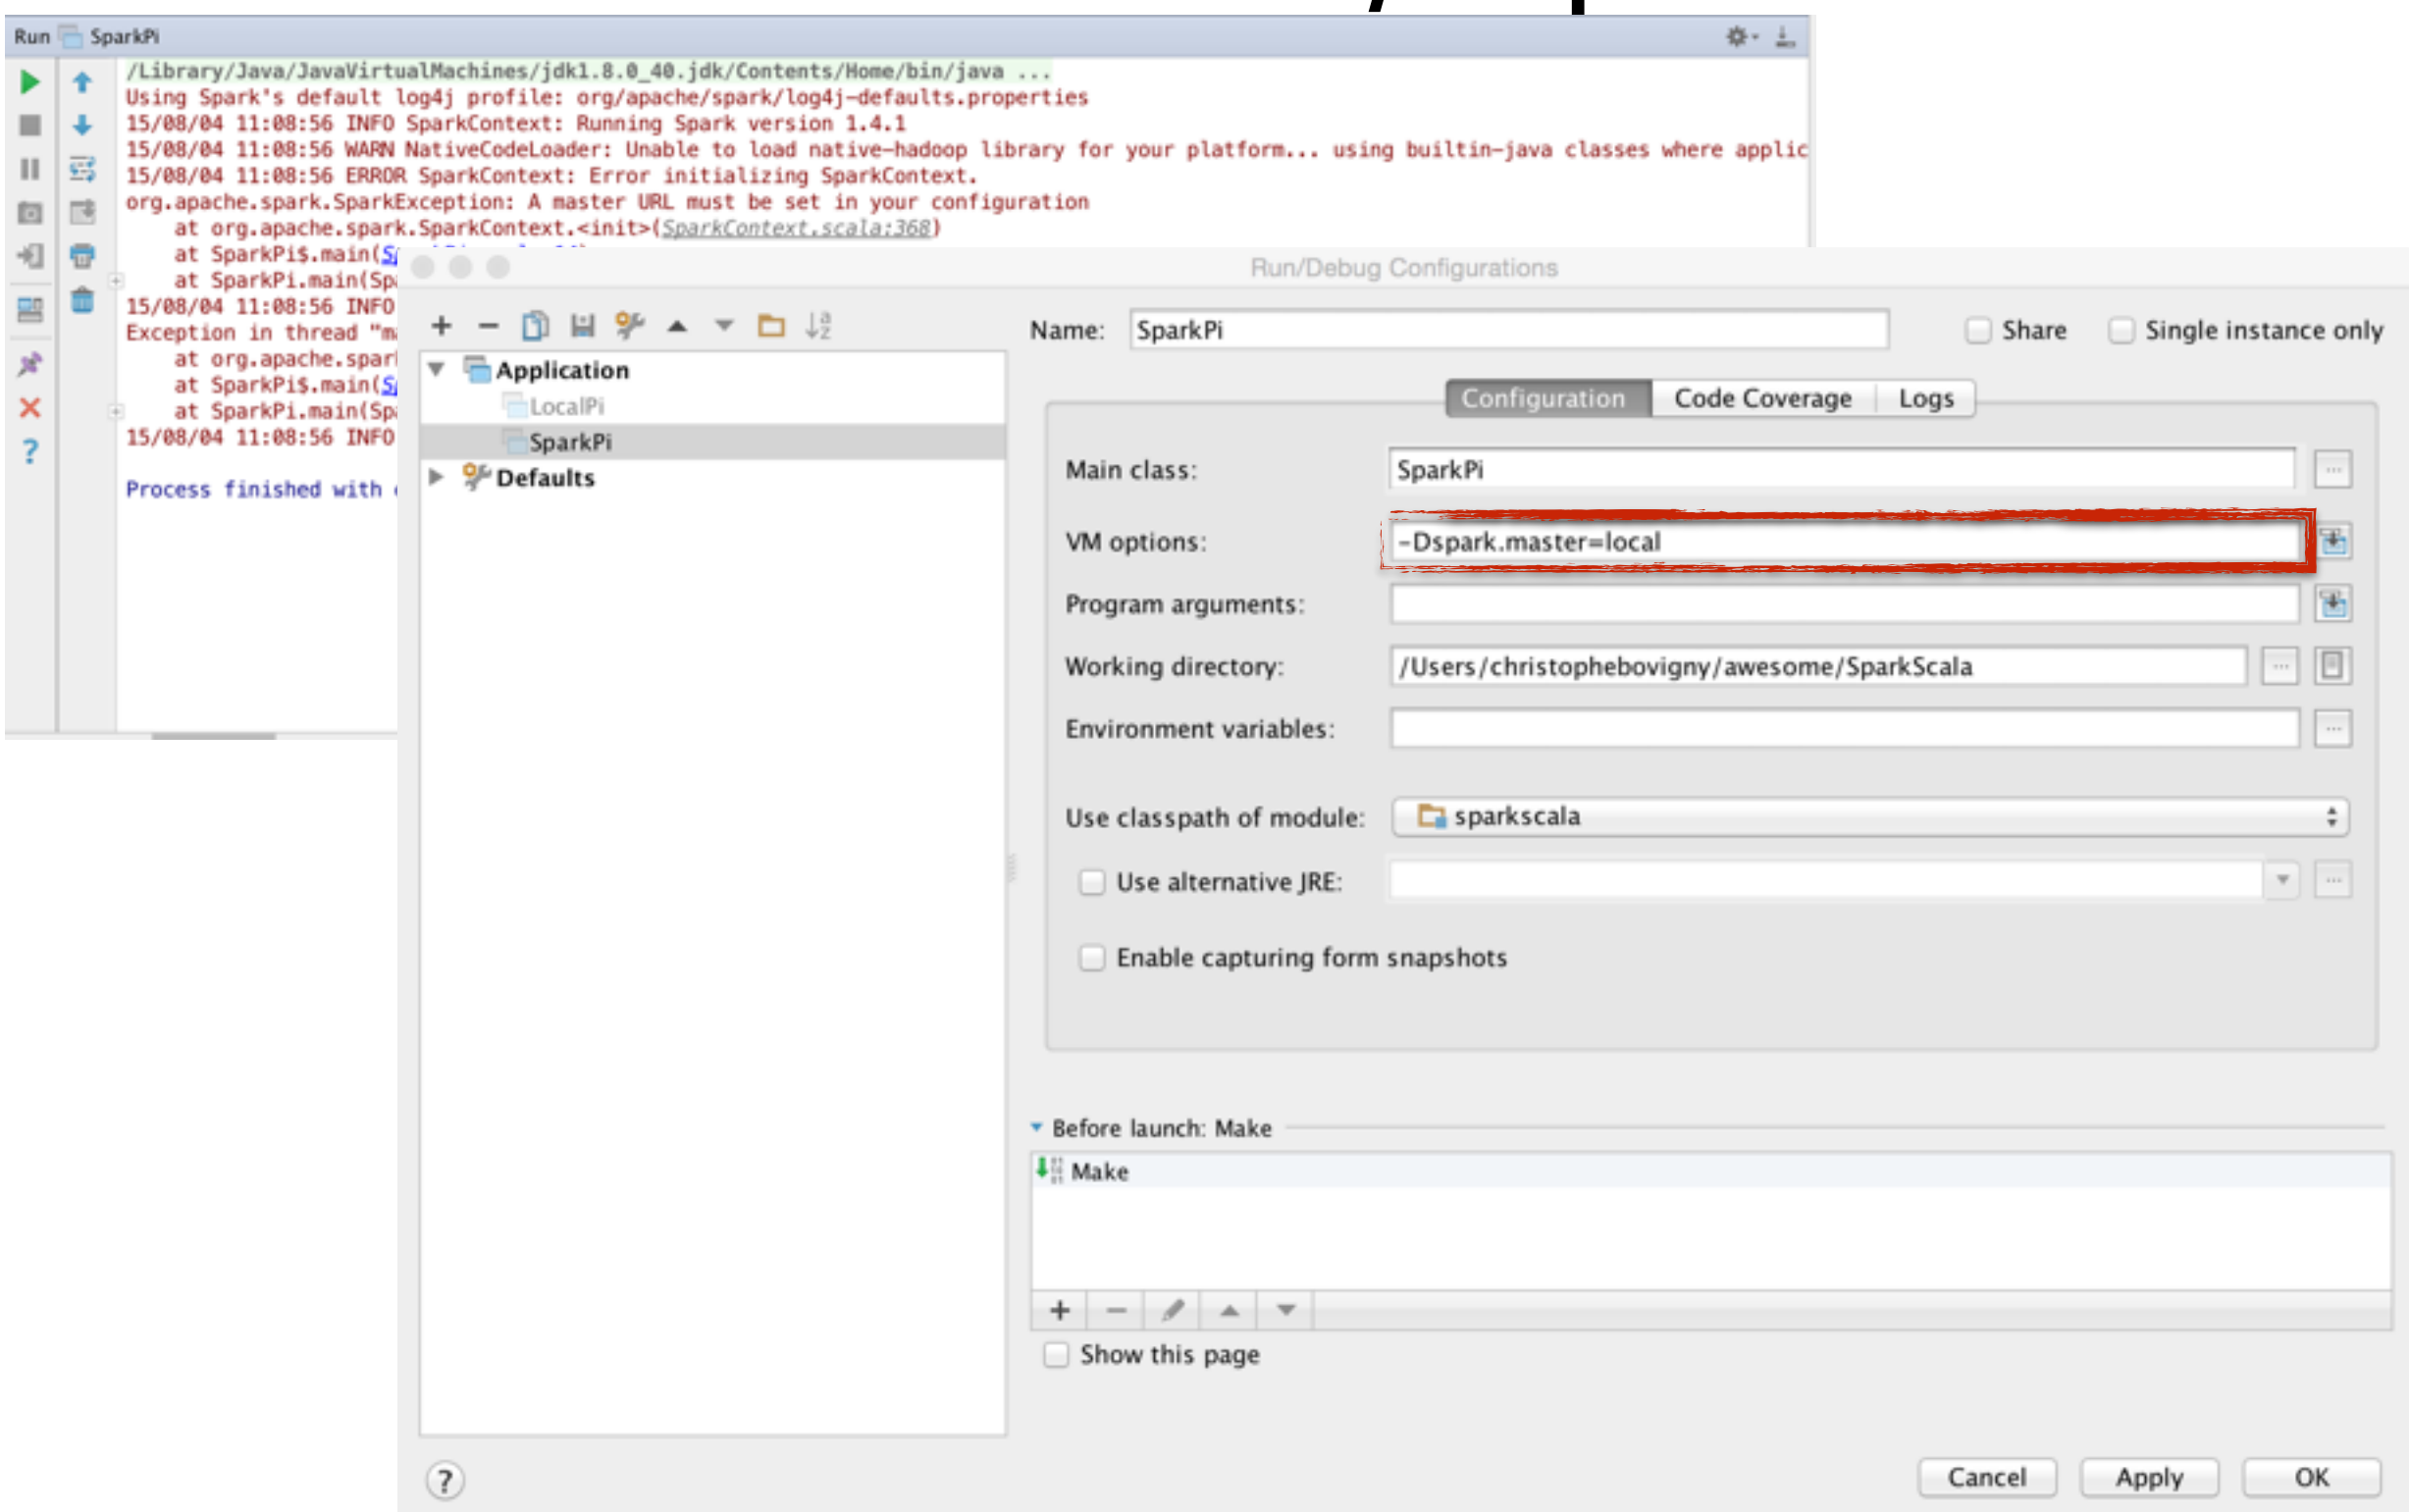
Pi / Local



Pi / Spark



Run Pi /Spark

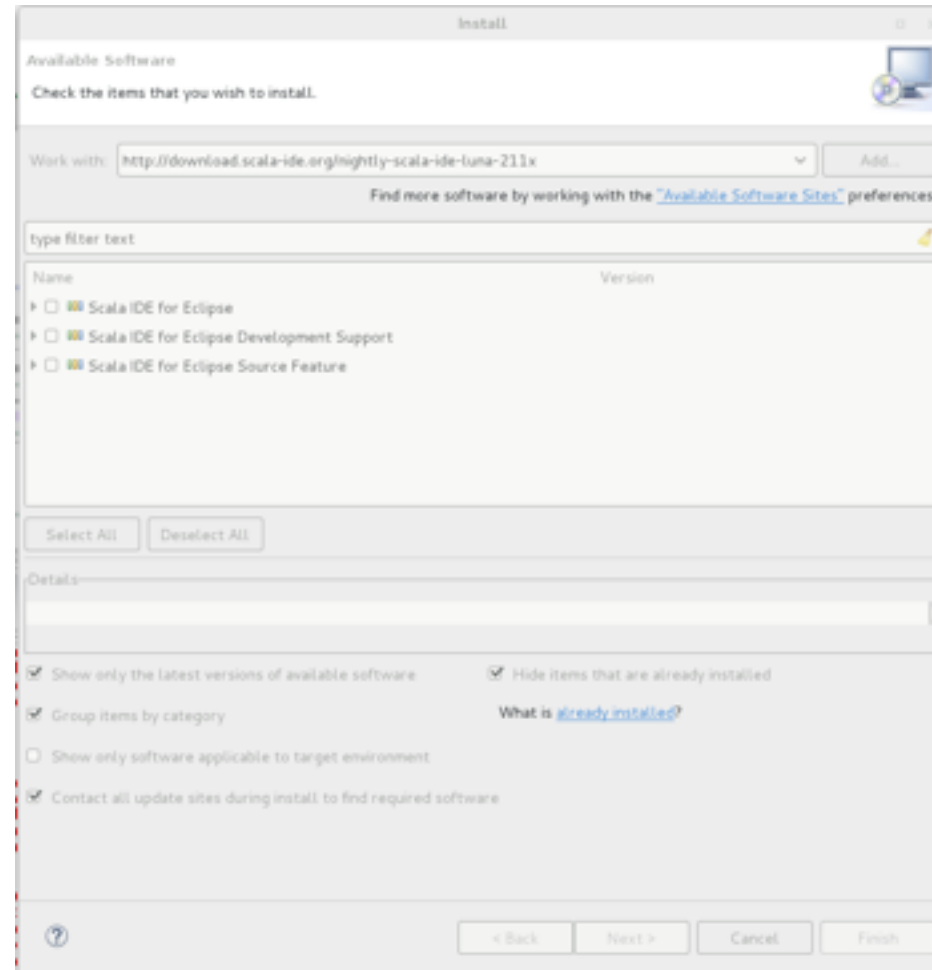


Setup a new Project

- Install sbt for your machine

www.scala-sbt.org/release/tutorial/Setup.html

- In eclipse Help -> Install New Software





Setup a new Project

- Create the project folder structure : *mkdir -p src/main/scala*
- Create a file in `src/main/scala` : *echo 'object Hi { def main(args: Array[String]) = println(« Hi !»)' > hw.scala*
- To allow managing dependencies, project name, Scala Version etc... create a file named *build.sbt* in the project root.

A screenshot of a code editor showing the content of a `build.sbt` file. The file is named `SparkPi.scala` and `build.sbt`. The code defines a project named `SparkScala` with version `1.0` and Scala version `2.11.7`. It also lists library dependencies for Spark, Kafka, and the Servlet API.

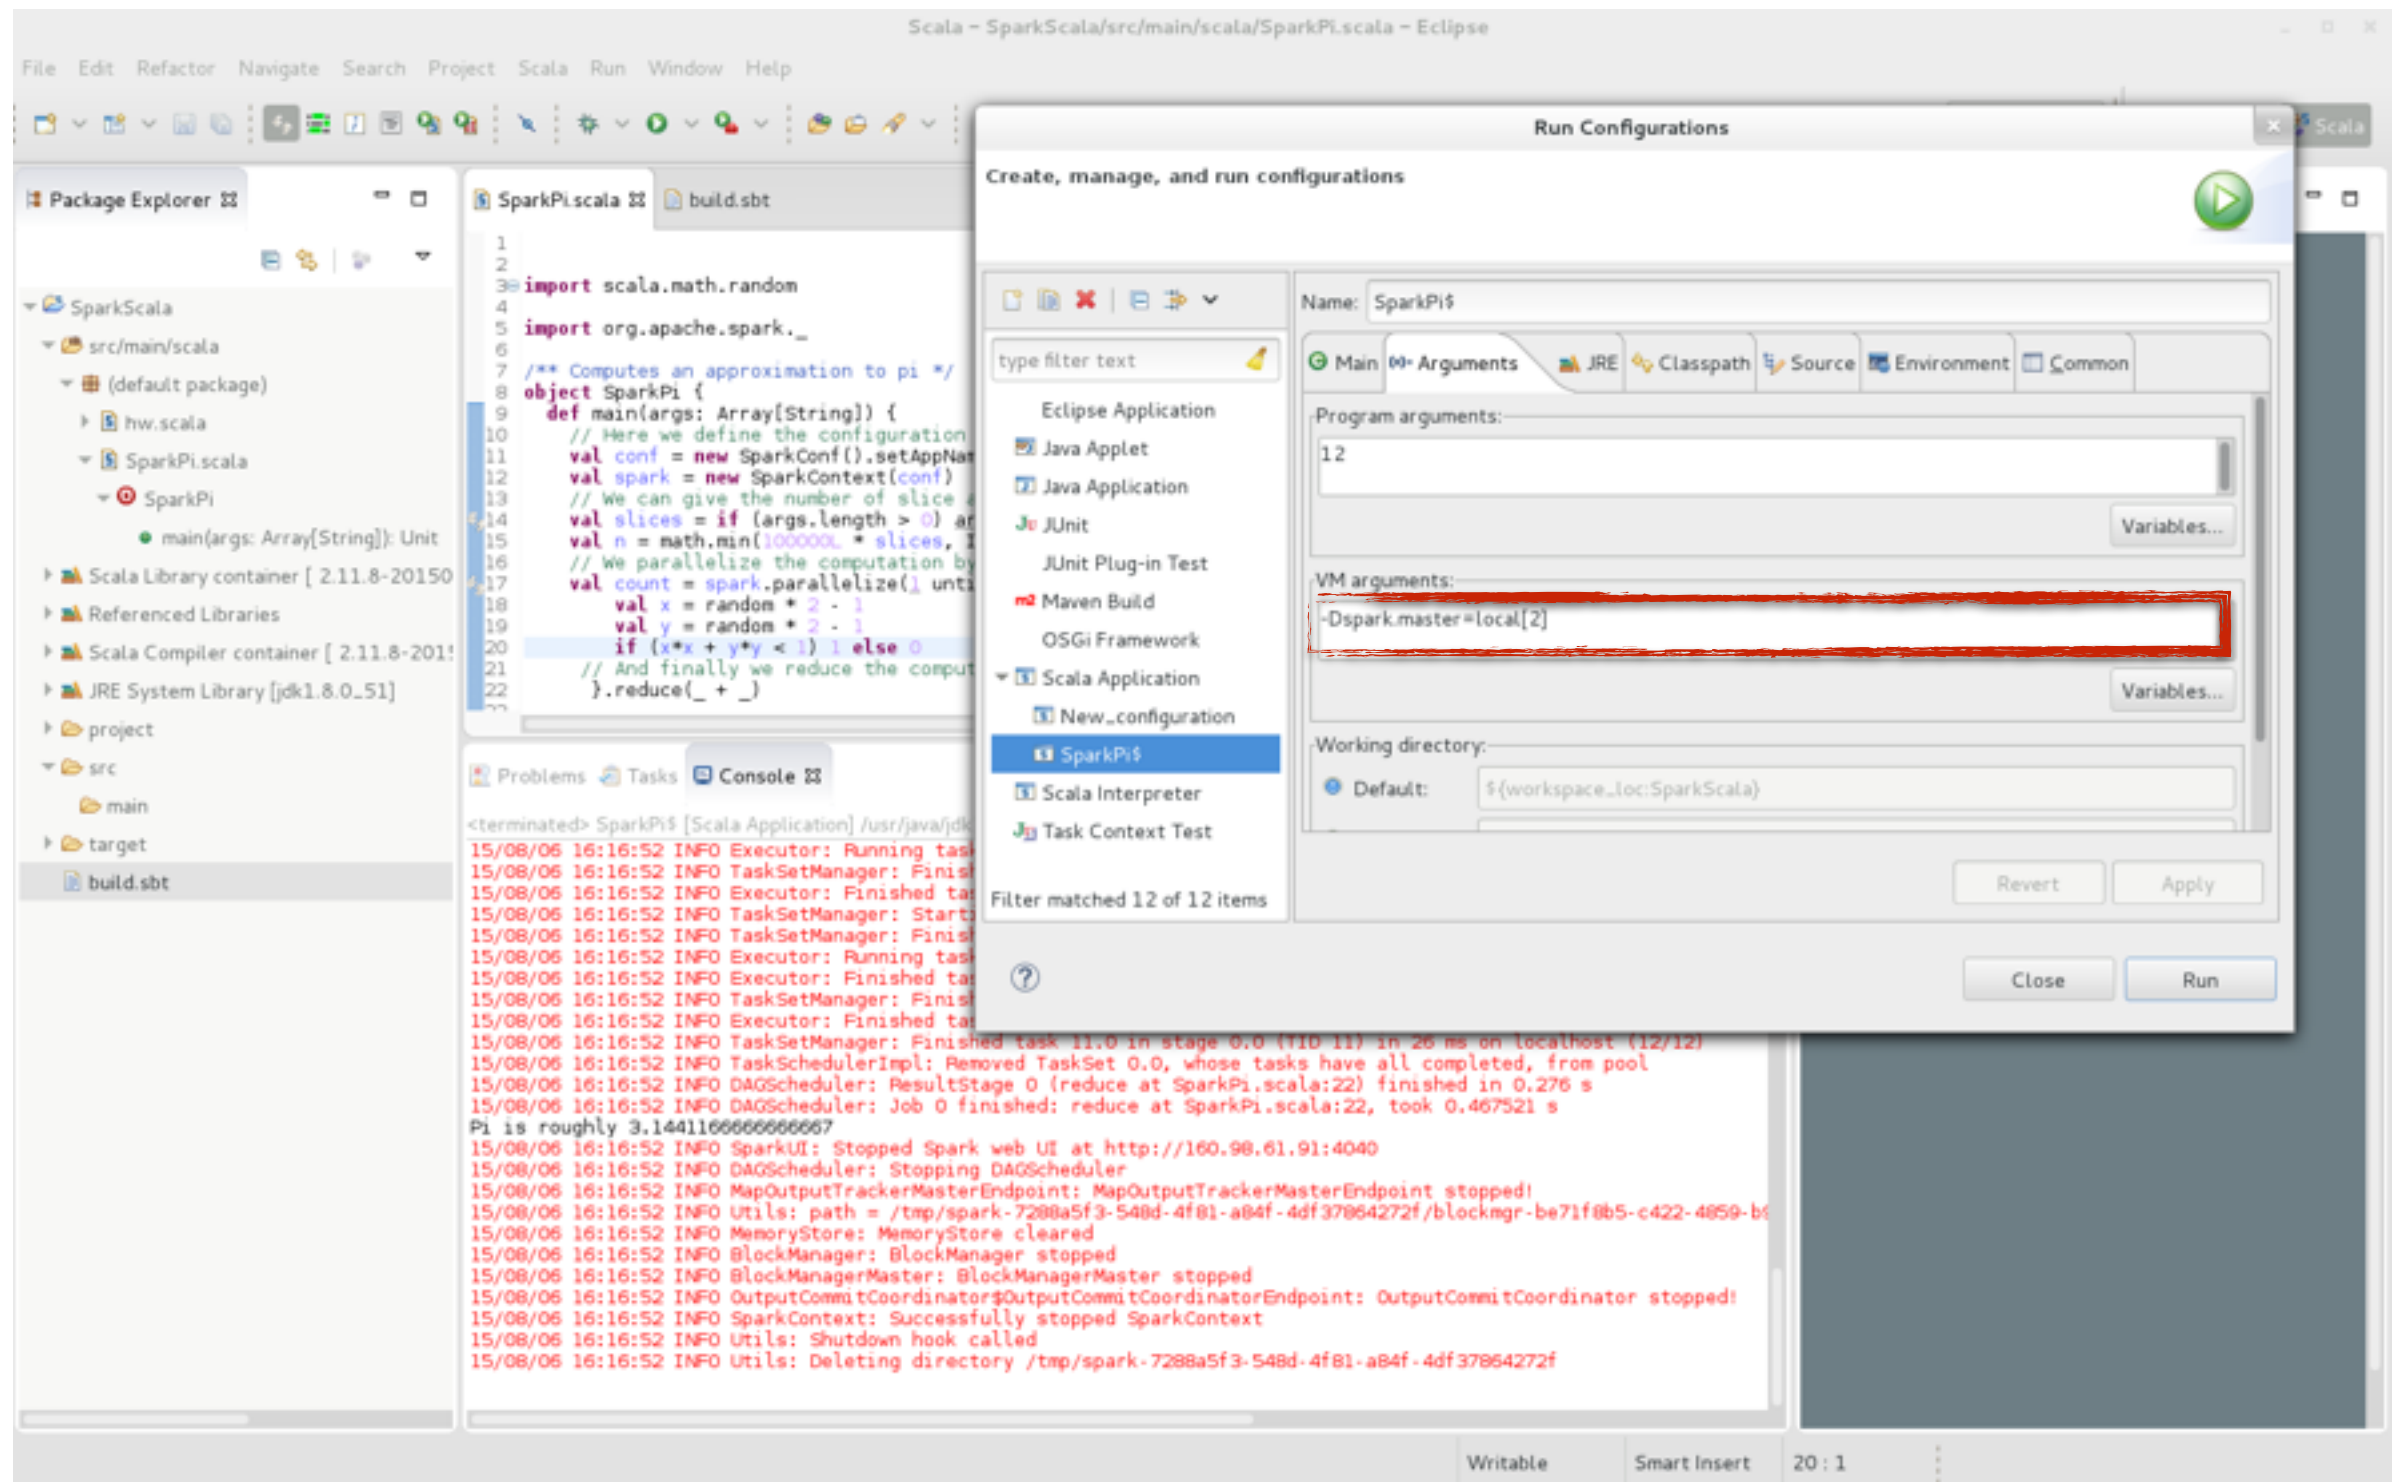
```
1 name := "SparkScala"
2
3 version := "1.0"
4
5 scalaVersion := "2.11.7"
6
7
8 libraryDependencies ++= Seq(
9   "org.apache.spark" %% "spark-streaming" % "1.4.1",
10  "org.apache.spark" %% "spark-streaming-kafka" % "1.4.1",
11  "javax.servlet" % "javax.servlet-api" % "3.0.1"
12)
13
14
```



Setup a new Project

- Open \$ ~/.sbt/0.13/plugins/build.sbt
*addSbtPlugin(«com.typesafe.sbteclipse » %
« sbteclipse-plugin » % 2.2.0)*
- In the project root folder : *sbt eclipse with-
source=true*
- File->Import-> Existing Projects into Workspace

Run Spark Project



Scala Spark Examples

- <https://github.com/apache/spark/tree/master/examples/src/main/scala/org/apache/spark/examples>
- <http://spark.apache.org>

